

## The System in One Flow

The product is a conversation that produces a ranked delegation plan. Three stages. The intake conversation surfaces five signals about how the owner spends their time. The prompt turns those signals into a categorized, ranked plan using the DRIP framework. The output is structured JSON that the frontend renders. Everything downstream keys off the schema, which is why we locked it first.

**Why?** This also defines our seam. I own the contract from questions to JSON output, which is the prompt, categorization, ranking, and schema. You wire that contract into the backend and database. We meet at the schema.

## The Question Set (Final for v1)

Five signals, plus a short opener for business context and team. Every question maps to exactly one signal, so nothing is arbitrary.

**Why?** A question earns its place only if it feeds a downstream step. If it feeds nothing, we cut it. If a step has no question feeding it, that is a gap. This is the test for the whole intake.

- **Business context (new).** Tell me about your business: what you do, team size, stage, industry.
  - Feeds: framing.
  - **Why?** The model needs to know who it is talking to before it can personalize, and this powers the adaptive follow-ups. Added because Amanda, Nichole, and Ashley all flagged that without context up front the recommendations come out generic.
- **Existing team (new, moved up).** Who is already on your team, even part-time or contract, and what does each person own? Optional org chart upload.
  - Feeds: constraints, and acts as a guardrail.
  - **Why?** This is the fix for the prototype bug where it recommended delegating marketing the team already owned. Ashley wants it moved up because knowing the team changes the entire who-takes-it answer. It is also what makes the output consultative instead of just go hire a VA.
- **Tasks, the week (unchanged).** Walk me through a typical week.

Feeds: tasks.

  - **Why?** This is extraction. Owners do not know what to delegate because it is all in their head, so we pull it out through conversation, not a form.
  - Monday - Thursday or just list out the tasks
- **Time per task (unchanged).** Roughly how many hours a week does each take?

Feeds: time.

  - **Why?** Ranking fuel and the headline number, hours reclaimed. Good candidate for a grid widget so people do not have to type, which Amanda and Nichole both asked for.

- Maybe add an extra column
- **Energy (unchanged).** Which of these drain you, and which give you energy?
  - Feeds: energy, DRIP axis one.
    - From the widget you can assume and then place them - in different quadrants
    - Visual and colorful (1-5) scale
- **Value (unchanged).** Which can only you do well, and which could someone else handle?
  - Feeds: value, DRIP axis two.
- **History and money (edited per feedback).** Have you delegated before, what worked, and what And what is your take-home or target income?
- got in the way? Feeds: constraints, plus the rate math.
  - **Why?** Two edits here. Nichole wanted the delegation question to ask about successes and roadblocks, not a yes or no.
  - **Ashley was explicit that we ask personal take-home or target income, not company revenue, because revenue made the rate math come out too high to be believable.**

**Adaptive, not longer. The backbone is the required ground. The model asks follow-ups only where they matter.**

### The Output Schema

The JSON shape does not change, because your backend builds to it. Here it is, with the two changes from the latest feedback marked inline.

```
{
  "summary": "string",
  "total_time_saved_hours_week": 0,
  "plan": [
    {
      "rank": 1,
      "task": "string",
      "drip_quadrant": "Delegation | Replacement | Investment | Production",
      "reasoning": "string",
      "who_takes_it": "string <-- CHANGE 1: now route + role",
      "                "      e.g. Existing team member: Ops Manager",
      "                "      or System, or External hire",
      "time_saved_hours_week": 0,
      "est_cost_to_delegate": "string",
      "net_recommendation": "string"
    }
  ],
}
```

```
"elevate_tier": "string <-- CHANGE 2: NEW, pending Ashley"
}
```

One ranked array so the frontend just maps over it. drip\_quadrant is a fixed set of values so the UI can group by quadrant and the eval set can check correctness. reasoning is its own field because the product is people-first and the owner needs the why, and because making the model justify its call improves accuracy. time and cost are separate

### **The two changes, flagged because they touch your backend**

1. **who\_takes\_it now expresses the route, not just a role.** The feedback was unanimous: do not assume hire a VA. First an existing team member if one fits, then a system, then outside help or Elevate. So the field carries both the route and the specific role or person.
2. **A new elevate\_tier field.** Ashley's name's the tier idea: the output should tell people where they land with Elevate.

### **What the Team Feedback Changes**

Most of it asks for things we already designed this week, which is the good. Here is who owns each piece.

#### **Mine to build**

- Consultative who-takes-it routing: existing person, then system, then external.
- The two question wording edits: personal income, and success and roadblock framing.
- Helper text or examples on the allocation table so people do not get stuck.

#### **Yours to wire**

- The two schema additions above: the route inside who\_takes\_it, and the elevate\_tier field once Ashley defines it.

### **Elevate to decide (we build to placeholders until they answer)**

- Tier names and which result maps to which. You already asked Ashley to clarify this.
- The voice. Ashley says it reads too SaaS. Our playbook Section 9 collects this.
- Moving off the borrowed buy back language onto Elevate's own framework.
- What happens after the start your handoff button.

### **What We Need to Decide Tomorrow**

Each of these should leave the meeting with a decision, an owner, or a date.

3. Freeze the v1 question set as written above. Owner: both.
4. Confirm the two schema additions. Owner: both.
7. Settle the right number of questions, your open item from Slack: the five signals plus the opener, with adaptive follow-ups on top. Confirm we are aligned.

- Calendar and to-do connectors: deferred to v2. They feed only the tasks signal and miss energy and value, and each is a real OAuth build. The conversation is the moat, so we do not automate it away for v1.
- Org chart upload: v1.5 fast-follow behind the text capture. It supplements the conversation and the model can read an image, but most small owners do not have one, and a chart shows reporting lines rather than who owns what. So text stays the primary path.
- RAG: v2, after prompt engineering is proven in v1.